

Data loading

We are combining the three datasets in this step with name, review ratings, review, and the number of people that found the review to be helpful

```
In [13]: import pandas as pd

# Replace paths with your exact locations if needed
fn1 = "Datafiniti_Amazon_Consumer_Reviews_of_Amazon_Products_May19.csv"
fn2 = "1429_1.csv"
fn3 = "Datafiniti_Amazon_Consumer_Reviews_of_Amazon_Products.csv"

cols = ["name", "reviews.rating", "reviews.text", "reviews.numHelpful"]

# Load each, ignoring missing useful outer columns
df1 = pd.read_csv(fn1, usecols=cols)
df2 = pd.read_csv(fn2, usecols=cols)
df3 = pd.read_csv(fn3, usecols=cols)

df = pd.concat([df1, df2, df3], ignore_index=True)
df.to_csv("combined_reviews.csv", index=False)

print("✅ Combined DataFrame saved to 'combined_reviews.csv'")
print("Shape:", df.shape)
print("NaNs in helpful:", df["reviews.numHelpful"].isna().sum())
```

<ipython-input-13-1568624014>:13: DtypeWarning: Columns (1) have mixed type s. Specify dtype option on import or set low_memory=False.

```
df2 = pd.read_csv(fn2, usecols=cols)
✅ Combined DataFrame saved to 'combined_reviews.csv'
Shape: (67992, 4)
NaNs in helpful: 12746
```

Data Preprocessing

Sentiment Classification

We are classifying based of the review rating of 3 and lower to be negative and 4 and above to be positive reviews.

```
In [16]: import pandas as pd

# Load combined CSV
df = pd.read_csv("combined_reviews.csv")

# Drop missing review text
df = df.dropna(subset=["reviews.text"])
```

```
# Reclassify using your rule: ≤ 3 = negative, ≥ 4 = positive
df["sentiment"] = df["reviews.rating"].apply(
    lambda r: "negative" if r <= 3 else "positive"
)
df.to_csv("labeled_reviews.csv", index=False)
# Show counts before balancing
counts_before = df["sentiment"].value_counts()
print("Counts before balancing:\n", counts_before)
```

Counts before balancing:

```
sentiment
positive    62579
negative     5412
Name: count, dtype: int64
```

Extract top helpful reviews

For later, it will be useful to know what were the most helpful reviews based on the reviews.numHelpful column for each product

```
In [17]: import pandas as pd

# Load the labeled dataset
df = pd.read_csv("labeled_reviews.csv")

# Replace NaNs in helpful counts with 0
df["reviews.numHelpful"] = df["reviews.numHelpful"].fillna(0)

# Define how many top reviews to pick per class per product
TOP_N = 5

# Sort and group to pick
top_reviews = (
    df
    .sort_values(by="reviews.numHelpful", ascending=False)
    .groupby(["name", "sentiment"])
    .head(TOP_N)
    .reset_index(drop=True)
)

# Quick check
print("Sample top helpful reviews:")
print(top_reviews[["name", "sentiment", "reviews.numHelpful"]].head(10))
print("\nCounts per sentiment per product (first few):")
print(top_reviews.groupby(["sentiment", "name"]).size().unstack(fill_value=0))

# Save for LLM processing
top_reviews.to_csv("top_helpful_reviews.csv", index=False)
print(f"\n✅ Top {TOP_N} helpful reviews per sentiment per product saved to")
```

Sample top helpful reviews:

	name	sentiment	\
0	Fire Tablet, 7 Display, Wi-Fi, 8 GB – Includes...	positive	
1	Fire Tablet, 7 Display, Wi-Fi, 8 GB – Includes...	positive	
2	Amazon Kindle Lighted Leather Cover,,, \r\nAmaz...	negative	
3	Fire Tablet, 7 Display, Wi-Fi, 8 GB – Includes...	positive	
4	Oem Amazon Kindle Power Usb Adapter Wall Trave...	negative	
5	AmazonBasics Bluetooth Keyboard for Android De...	positive	
6	Amazon Tap Smart Assistant Alexaenabled (black...	positive	
7	Fire Kids Edition Tablet, 7 Display, Wi-Fi, 16...	positive	
8	AmazonBasics Bluetooth Keyboard for Android De...	positive	
9	Echo (White),,,, \r\nEcho (White),,,,	negative	

	reviews.numHelpful
0	780.0
1	740.0
2	730.0
3	650.0
4	621.0
5	525.0
6	434.0
7	355.0
8	345.0
9	292.0

Counts per sentiment per product (first few):

name	All-New Fire 7 Tablet with Alexa, 7" Display, 8 GB – Marine Blue	\
sentiment		
negative		4
positive		5

name	All-New Fire HD 8 Kids Edition Tablet, 8 HD Display, 32 GB, Blue Kid-Proof Case	\
sentiment		
negative		5
positive		5

name	All-New Fire HD 8 Kids Edition Tablet, 8 HD Display, 32 GB, Pink Kid-Proof Case	\
sentiment		
negative		5
positive		5

name	All-New Fire HD 8 Tablet with Alexa, 8 HD Display, 16 GB, Marine Blue – with Special Offers	\
sentiment		
negative		5
positive		5

name	All-New Fire HD 8 Tablet with Alexa, 8 HD Display, 32 GB, Marine Blue – with Special Offers	\
sentiment		
negative		5
positive		5

✓ Top 5 helpful reviews per sentiment per product saved to 'top_helpful_reviews.csv'

Train Sentiment Classification Model

Data prep

Using the labeled data we can:

1. Map negative and positive reviews
2. Run a train-test split
3. Perform Tokenization using DistilBert
4. Save the encodings for Train and Test in JSON to further use

```
In [21]: import pandas as pd
from sklearn.model_selection import train_test_split
from transformers import DistilBertTokenizerFast
import json

# Load labeled data
df = pd.read_csv("labeled_reviews.csv").dropna(subset=["reviews.text"])
label_map = {"negative": 0, "positive": 1}
df["label"] = df["sentiment"].map(label_map)

# Split into train/test
train_df, test_df = train_test_split(df, test_size=0.2, stratify=df["label"])

# Tokenization
tokenizer = DistilBertTokenizerFast.from_pretrained("distilbert-base-uncased")
train_enc = tokenizer(train_df["reviews.text"].tolist(), truncation=True, padding="max_length")
test_enc = tokenizer(test_df["reviews.text"].tolist(), truncation=True, padding="max_length")

# Convert BatchEncoding to plain dicts
train_enc_dict = {k: v for k, v in train_enc.items()}
test_enc_dict = {k: v for k, v in test_enc.items()}

# Save as JSON
with open("train_enc.json", "w") as f:
    json.dump({"encodings": train_enc_dict, "labels": train_df["label"].tolist(), "train": True}, f)

with open("test_enc.json", "w") as f:
    json.dump({"encodings": test_enc_dict, "labels": test_df["label"].tolist(), "train": False}, f)

# Save tokenizer and confirm
tokenizer.save_pretrained("tokenizer")
print("✓ Tokenized encodings and labels saved as JSON. Let's load and verify next.")
```

✓ Tokenized encodings and labels saved as JSON. Let's load and verify next.

Train with Weighted Loss

Load the JSON encodings for train and test and instantiate the datasets to be used for sentiment analysis

```
In [22]: import json
import torch
from torch.utils.data import Dataset

# Load JSON data
with open("train_enc.json") as f:
    train_data = json.load(f)
with open("test_enc.json") as f:
    test_data = json.load(f)

# Define Dataset class
class SentimentDataset(Dataset):
    def __init__(self, encodings, labels):
        self.encodings = encodings
        self.labels = labels

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, idx):
        item = {k: torch.tensor(v[idx]) for k, v in self.encodings.items()}
        item["labels"] = torch.tensor(self.labels[idx])
        return item

# Instantiate datasets
train_ds = SentimentDataset(train_data["encodings"], train_data["labels"])
test_ds = SentimentDataset(test_data["encodings"], test_data["labels"])
print(f"✅ Loaded datasets: {len(train_ds)} train, {len(test_ds)} test samples")
```

✅ Loaded datasets: 54392 train, 13599 test samples.

Doing the training for the sentiment analysis classifier using DistilBert and accounting for imbalances using weighted loss we are overriding the compute_loss method in the trainer to plug in class weights. So it will pay attention to the under represented class, which is the negative reviews

```
In [27]: import torch
import torch.nn as nn
from transformers import DistilBertForSequenceClassification, Trainer, TrainingArguments

import pandas as pd
df = pd.read_csv("labeled_reviews.csv")
counts = df["sentiment"].value_counts()
total = len(df)
class_weights = torch.tensor(
    [total/counts["negative"], total/counts["positive"]],
    dtype=torch.float)

class WeightedTrainer(Trainer):
    #overriding here
    def compute_loss(self, model, inputs, return_outputs=False, **kwargs):
        labels = inputs.pop("labels")
        outputs = model(**inputs)
```

```

        logits = outputs.logits
        loss_fn = nn.CrossEntropyLoss(weight=class_weights.to(logits.device))
        loss = loss_fn(logits, labels)
        return (loss, outputs) if return_outputs else loss

training_args = TrainingArguments(
    output_dir="sentiment_model",
    num_train_epochs=3,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    weight_decay=0.01,
    logging_steps=100,
    save_total_limit=2,
    do_train=True,
    do_eval=True,
    eval_steps=500,
    save_steps=500
)

model = DistilBertForSequenceClassification.from_pretrained("distilbert-base-uncased")
trainer = WeightedTrainer(
    model=model,
    args=training_args,
    train_dataset=train_ds,
    eval_dataset=test_ds
)

#Train
trainer.train()
trainer.save_model("sentiment_model")
print("✅ Model trained and saved successfully.")

```

Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight', 'pre_classifier.bias', 'pre_classifier.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

 [10200/10200 30:59, Epoch 3/3]

Step	Training Loss
100	0.526500
200	0.537200
300	0.520000
400	0.452500
500	0.507000
600	0.462800
700	0.509400
800	0.458900
900	0.484300
1000	0.417000
1100	0.528300
1200	0.496000
1300	0.432900
1400	0.571300
1500	0.424000
1600	0.429100
1700	0.467600
1800	0.444600
1900	0.490200
2000	0.478200
2100	0.456800
2200	0.506000
2300	0.420400
2400	0.427100
2500	0.483100
2600	0.433800
2700	0.369200
2800	0.431400
2900	0.437100
3000	0.353800
3100	0.478100

Step	Training Loss
3200	0.417000
3300	0.428600
3400	0.418800
3500	0.359300
3600	0.319900
3700	0.303600
3800	0.313600
3900	0.291100
4000	0.341100
4100	0.365100
4200	0.283300
4300	0.285800
4400	0.278200
4500	0.266500
4600	0.296500
4700	0.330900
4800	0.371900
4900	0.368600
5000	0.320000
5100	0.314700
5200	0.328300
5300	0.340800
5400	0.259100
5500	0.333400
5600	0.288200
5700	0.278700
5800	0.338200
5900	0.280000
6000	0.260800
6100	0.361700
6200	0.285600
6300	0.292500

Step	Training Loss
6400	0.379300
6500	0.377000
6600	0.329800
6700	0.339200
6800	0.295200
6900	0.184200
7000	0.227000
7100	0.209600
7200	0.140200
7300	0.219700
7400	0.106400
7500	0.070500
7600	0.198200
7700	0.216200
7800	0.199100
7900	0.226600
8000	0.141100
8100	0.170100
8200	0.146000
8300	0.199900
8400	0.180700
8500	0.305700
8600	0.122800
8700	0.183700
8800	0.156000
8900	0.107800
9000	0.140800
9100	0.134900
9200	0.166000
9300	0.150700
9400	0.184000
9500	0.217100

Step	Training Loss
9600	0.236600
9700	0.157100
9800	0.183500
9900	0.208300
10000	0.271300
10100	0.137200
10200	0.165600

✅ Model trained and saved successfully.

```
In [29]: from sklearn.metrics import classification_report
import json

results = trainer.predict(test_ds)
y_pred = results.predictions.argmax(axis=1)
true_labels = [item["labels"].item() for item in test_ds]

report = classification_report(
    true_labels,
    y_pred,
    target_names=["negative", "positive"],
    output_dict=True
)

with open("classification_report.json", "w") as f:
    json.dump(report, f, indent=4)

print("✅ Evaluation results:\n", json.dumps(report, indent=2))
```

✅ Evaluation results:

```
{
  "negative": {
    "precision": 0.8391959798994975,
    "recall": 0.7717190388170055,
    "f1-score": 0.8040442946557534,
    "support": 1082.0
  },
  "positive": {
    "precision": 0.9804030466518565,
    "recall": 0.9872173843572741,
    "f1-score": 0.9837984156681661,
    "support": 12517.0
  },
  "accuracy": 0.9700713287741746,
  "macro avg": {
    "precision": 0.909799513275677,
    "recall": 0.8794682115871398,
    "f1-score": 0.8939213551619598,
    "support": 13599.0
  },
  "weighted avg": {
    "precision": 0.9691679524371309,
    "recall": 0.9700713287741746,
    "f1-score": 0.9694963376524716,
    "support": 13599.0
  }
}
```

Negative: Precision (83.9%): When the model predicts "negative," it's correct 83.9% of the time.

Recall (77.2%): The model correctly identifies 77.2% of actual negative reviews.

F1-score (80.4%): Harmonic mean of precision and recall, giving a balanced performance metric.

Positive: Precision (98.0%): Very high confidence that predicted positives are actually positive.

Recall (98.7%): Nearly all actual positive reviews were detected.

F1-score (98.4%): Shows excellent performance.

I think this is really good considering the heavy imbalance between positive and negative reviews

Inference & Product-Level Sentiment Analysis

To predict sentiment labels for every review in the dataset using we trained DistilBERT model, and now we aggregate those results to summarize sentiment per product.

Loading trained sentiment classifier model.

Tokenizing all product reviews in batches.

Running inference efficiently on GPU.

Saving predictions alongside the original data.

```
In [ ]: import pandas as pd

df = pd.read_csv("combined_reviews.csv")
df = df[df["reviews.text"].apply(lambda x: isinstance(x, str))].copy()
```

```
In [48]: from transformers import DistilBertTokenizerFast
import torch

tokenizer = DistilBertTokenizerFast.from_pretrained("tokenizer")
texts = df["reviews.text"].tolist()

input_ids_list, attention_masks_list = [], []
batch_size = 1000

# This loop ensures each batch of 1000 reviews is:
# Tokenized into input_ids and attention_mask
# Padded and truncated to 256 tokens

for i in range(0, len(texts), batch_size):
    batch = texts[i : i + batch_size]
    enc = tokenizer(
        batch,
        truncation=True,
        padding="max_length",
        max_length=256,
        return_tensors="pt"
    )
    input_ids_list.append(enc["input_ids"])
    attention_masks_list.append(enc["attention_mask"])

# Then everything is concatenated into tensors:
input_ids = torch.cat(input_ids_list)
attention_masks = torch.cat(attention_masks_list)
```

```
In [51]: import time
import torch
from transformers import DistilBertForSequenceClassification
from torch.utils.data import DataLoader, TensorDataset
from torch.cuda.amp import autocast
```

```

# Load model and tokenizer
model = DistilBertForSequenceClassification.from_pretrained("sentiment_model")

#Setup Inference Pipeline
dataset = TensorDataset(input_ids, attention_masks)
loader = DataLoader(dataset, batch_size=128)

#Running the inference pipeline
start = time.time()
preds = []
with torch.no_grad():
    for ids, masks in loader:
        ids, masks = ids.cuda(), masks.cuda()
        with autocast():
            outputs = model(input_ids=ids, attention_mask=masks)
            preds.extend(torch.argmax(outputs.logits, dim=1).tolist())
end = time.time()

print(f"✅ Inference complete in {end - start:.2f} seconds.")
df["pred_label"] = preds
df["pred_sentiment"] = df["pred_label"].map({0: "negative", 1: "positive"})
# Attach predictions to df
df.to_csv("all_reviews_with_preds.csv", index=False)

```

```

<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
<ipython-input-51-3415986705>:18: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with autocast():
    ✅ Inference complete in 29.83 seconds.

```

Product level summaries

```
In [52]: import pandas as pd

df = pd.read_csv("all_reviews_with_preds.csv")
# Aggregate Sentiment Stats by Product
stats = (
    df.groupby("name")["pred_sentiment"]
      .value_counts(normalize=True)
      .unstack(fill_value=0)
)

#Formatting to later use with LLM
stats["review_count"] = df.groupby("name").size()
stats.reset_index(inplace=True)
stats.rename(columns={"negative":"neg_pct", "positive":"pos_pct"}, inplace=True)

stats.to_csv("product_sentiment_stats.csv", index=False)
```

```
In [ ]: #Finding the top negative and positive based of number of helpful column

top_pos = df[df.pred_sentiment=="positive"].sort_values(["name","reviews.num
top_neg = df[df.pred_sentiment=="negative"].sort_values(["name","reviews.num

top_reviews = top_pos.merge(top_neg, on="name", suffixes=("_pos","_neg"))
top_reviews[["name","reviews.text_pos","reviews.text_neg"]].to_csv("top_revi
```

We now have

Quantitative stats: percent of positive/negative reviews per product + review counts

Qualitative review selection: Most helpful positive and negative review per product

These two together give us a strong base for:

Sentiment analysis

Product summaries

AI-generated overviews for shoppers or stakeholders

Saving Data

```
In [54]: # Package all project files into one ZIP
!zip -r aai510_full_project.zip \
    combined_reviews.csv \
    labeled_reviews.csv \
    all_reviews_with_preds.csv \
    product_sentiment_stats.csv \
    top_reviews_per_product.csv \
```

```
classification_report.json \  
top_helpful_reviews.csv \  
sentiment_model \  
tokenizer \  
train_enc.json \  
test_enc.json \  
train_ds.pt \  
test_ds.pt
```

```

adding: combined_reviews.csv (deflated 76%)
adding: labeled_reviews.csv (deflated 76%)
adding: all_reviews_with_preds.csv (deflated 76%)
adding: product_sentiment_stats.csv (deflated 73%)
adding: top_reviews_per_product.csv (deflated 72%)
adding: classification_report.json (deflated 63%)
adding: top_helpful_reviews.csv (deflated 77%)
adding: sentiment_model/ (stored 0%)
adding: sentiment_model/training_args.bin (deflated 52%)
adding: sentiment_model/model.safetensors (deflated 8%)
adding: sentiment_model/config.json (deflated 43%)
adding: sentiment_model/checkpoint-10200/ (stored 0%)
adding: sentiment_model/checkpoint-10200/training_args.bin (deflated 52%)
adding: sentiment_model/checkpoint-10200/model.safetensors (deflated 8%)
adding: sentiment_model/checkpoint-10200/config.json (deflated 43%)
adding: sentiment_model/checkpoint-10200/trainer_state.json (deflated 81%)
adding: sentiment_model/checkpoint-10200/rng_state.pth (deflated 25%)
adding: sentiment_model/checkpoint-10200/scheduler.pt (deflated 56%)
adding: sentiment_model/checkpoint-10200/optimizer.pt (deflated 30%)
adding: sentiment_model/runs/ (stored 0%)
adding: sentiment_model/runs/Jun15_20-43-37_a35f6563acc9/ (stored 0%)
adding: sentiment_model/runs/Jun15_20-43-37_a35f6563acc9/events.out.tfeven
ts.1750020219.a35f6563acc9.4851.3 (deflated 67%)
adding: sentiment_model/runs/Jun15_20-41-57_a35f6563acc9/ (stored 0%)
adding: sentiment_model/runs/Jun15_20-41-57_a35f6563acc9/events.out.tfeven
ts.1750020118.a35f6563acc9.4851.1 (deflated 61%)
adding: sentiment_model/runs/Jun15_20-42-26_a35f6563acc9/ (stored 0%)
adding: sentiment_model/runs/Jun15_20-42-26_a35f6563acc9/events.out.tfeven
ts.1750020148.a35f6563acc9.4851.2 (deflated 61%)
adding: sentiment_model/runs/Jun15_20-40-38_a35f6563acc9/ (stored 0%)
adding: sentiment_model/runs/Jun15_20-40-38_a35f6563acc9/events.out.tfeven
ts.1750020052.a35f6563acc9.4851.0 (deflated 61%)
adding: sentiment_model/checkpoint-10000/ (stored 0%)
adding: sentiment_model/checkpoint-10000/training_args.bin (deflated 52%)
adding: sentiment_model/checkpoint-10000/model.safetensors (deflated 8%)
adding: sentiment_model/checkpoint-10000/config.json (deflated 43%)
adding: sentiment_model/checkpoint-10000/trainer_state.json (deflated 81%)
adding: sentiment_model/checkpoint-10000/rng_state.pth (deflated 25%)
adding: sentiment_model/checkpoint-10000/scheduler.pt (deflated 56%)
adding: sentiment_model/checkpoint-10000/optimizer.pt (deflated 30%)
adding: tokenizer/ (stored 0%)
adding: tokenizer/special_tokens_map.json (deflated 42%)
adding: tokenizer/vocab.txt (deflated 53%)
adding: tokenizer/tokenizer.json (deflated 71%)
adding: tokenizer/tokenizer_config.json (deflated 75%)
adding: train_enc.json (deflated 96%)
adding: test_enc.json (deflated 96%)
adding: train_ds.pt (deflated 94%)
adding: test_ds.pt (deflated 94%)

```

```
In [55]: from google.colab import files
files.download('aai510_full_project.zip')
```

```
In [58]: from google.colab import files
```




```
files.download('product_sentiment_stats.csv')
files.download('top_reviews_per_product.csv')
files.download('all_reviews_with_preds.csv')
```

Visualizations

```
In [59]: stats = pd.read_csv("product_sentiment_stats.csv")
mean_pos = stats['pos_pct'].mean()
mean_neg = stats['neg_pct'].mean()

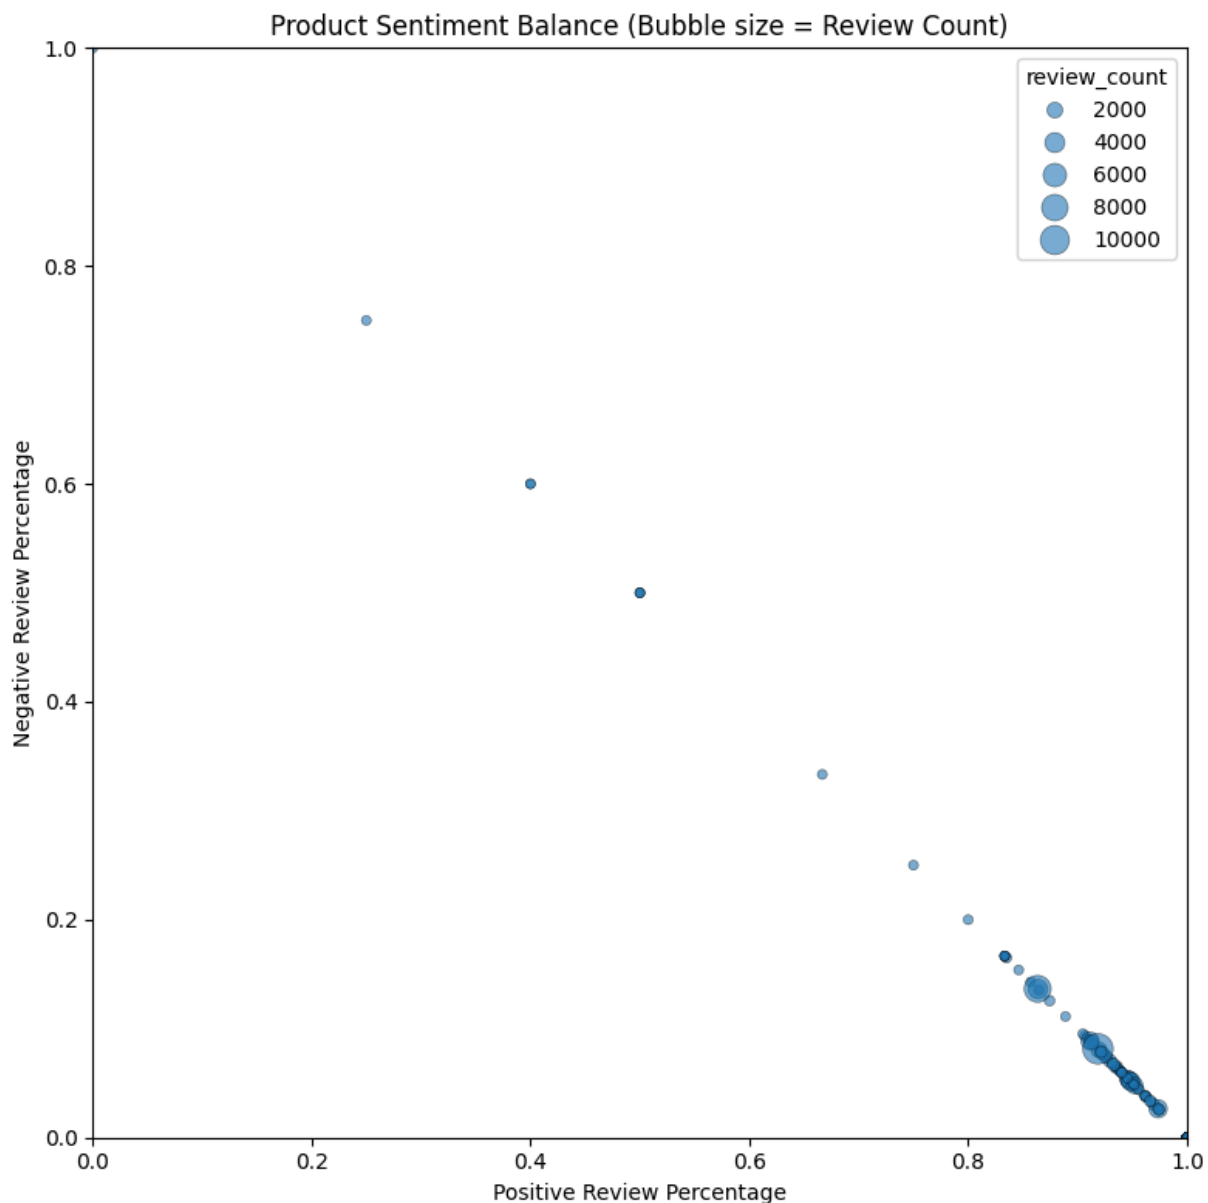
plt.figure(figsize=(6, 4))
plt.bar(['Mean Positive %', 'Mean Negative %'], [mean_pos, mean_neg], color=
plt.ylabel("Average Percentage Across Products")
plt.ylim(0, 1)
plt.tight_layout()
plt.show()
```



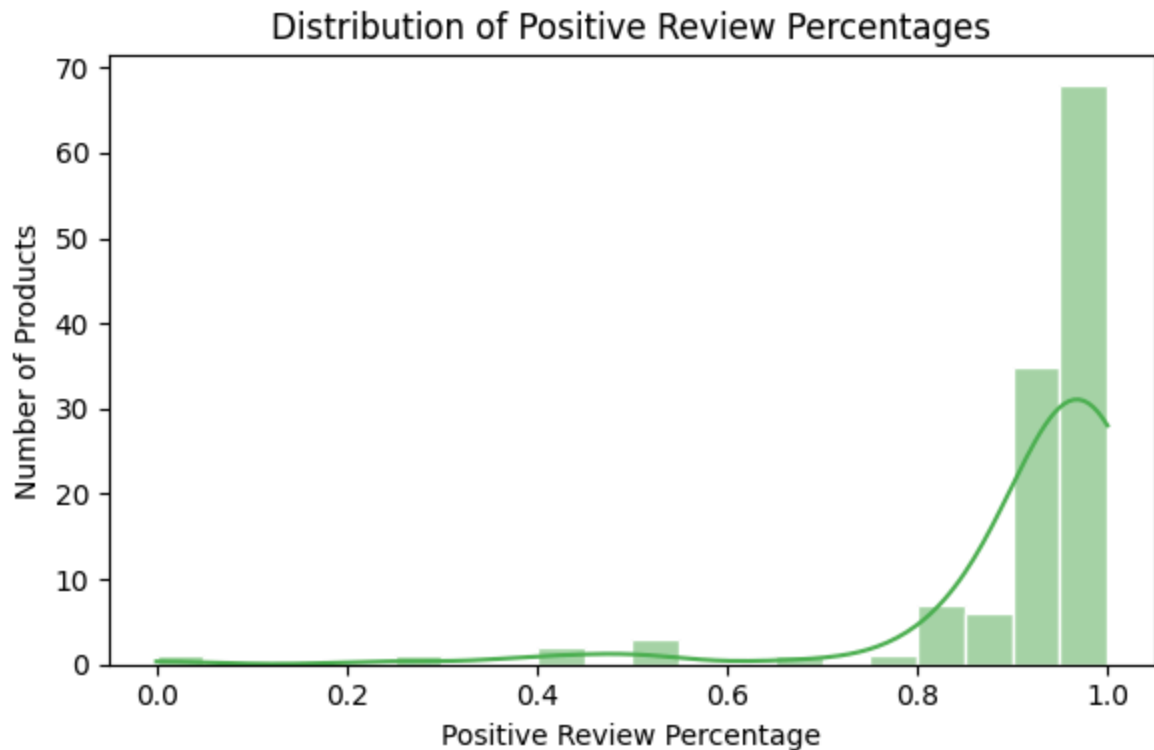
```
In [60]: import seaborn as sns

plt.figure(figsize=(8, 8))
sns.scatterplot(
    data=stats,
    x="pos_pct",
    y="neg_pct",
    size="review_count",
    sizes=(20, 200),
    alpha=0.6,
    edgecolor="k"
)
plt.title("Product Sentiment Balance (Bubble size = Review Count)")
```

```
plt.xlabel("Positive Review Percentage")
plt.ylabel("Negative Review Percentage")
plt.xlim(0, 1)
plt.ylim(0, 1)
plt.tight_layout()
```



```
In [61]: plt.figure(figsize=(6, 4))
sns.histplot(stats['pos_pct'], bins=20, kde=True, color="#4CAF50", edgecolor
plt.title("Distribution of Positive Review Percentages")
plt.xlabel("Positive Review Percentage")
plt.ylabel("Number of Products")
plt.tight_layout()
```




```
In [62]: top_neg = stats.sort_values('neg_pct', ascending=False).head(10)

plt.figure(figsize=(6, 5))
sns.barplot(
    data=top_neg,
    x='neg_pct',
    y='name',
    palette='Reds_r'
)
plt.title("Top 10 Products with Highest Negative Sentiment")
plt.xlabel("Negative Review Percentage")
plt.ylabel("Product Name")
plt.tight_layout()
```

```
<ipython-input-62-1267421246>:4: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
<ipython-input-62-1267421246>:13: UserWarning: Glyph 13 (
) missing from font(s) DejaVu Sans.
plt.tight_layout()
<ipython-input-62-1267421246>:13: UserWarning: Tight layout not applied. The
left and right margins cannot be made large enough to accommodate all Axes d
ecorations.
plt.tight_layout()
/usr/local/lib/python3.11/dist-packages/IPython/core/events.py:89: UserWarni
ng: Glyph 13 (
) missing from font(s) DejaVu Sans.
func(*args, **kwargs)
) missing from font(s) DejaVu Sans.
fig.canvas.print_figure(bytes_io, **kw)
```



```
In [63]: from wordcloud import WordCloud
```

```
top_reviews = pd.read_csv("top_reviews_per_product.csv")
text_pos = " ".join(top_reviews["reviews.text_pos"].astype(str).values)
text_neg = " ".join(top_reviews["reviews.text_neg"].astype(str).values)

wc_pos = WordCloud(width=400, height=200, background_color="white", colormap=
wc_neg = WordCloud(width=400, height=200, background_color="white", colormap=

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.imshow(wc_pos, interpolation='bilinear')
plt.axis("off")
plt.title("Positive Reviews Word Cloud")

plt.subplot(1, 2, 2)
plt.imshow(wc_neg, interpolation='bilinear')
plt.axis("off")
plt.title("Negative Reviews Word Cloud")
plt.tight_layout()
```



```
%%shell
jupyter nbconvert --to html /content/sentiment_analysis_my.ipynb
```

```
[NbConvertApp] Converting notebook /content/sentiment_analysis_my.ipynb to h
tml
```

```
[NbConvertApp] WARNING | Alternative text is missing on 5 image(s).
```

```
[NbConvertApp] Writing 1106909 bytes to /content/sentiment_analysis_my.html
```

Out[64]: